

Run-level collapse: same operator (effective N = exp(H)), two inputs

Δ effective topics (post-text embedding) Δ effective frames (LLM-judge dominant_frame)

