

Length-controlled Distinct-5 (lower means less lexical diversity)				
Condition	GPT-5 1 h	Gemini Flash Lite 1 h	GPT-5 5 h	Gemini Flash Lite 5 h
Empty	0.958 → 0.931	0.985 → 0.962	0.997 → 0.962	–
1 conspiracy	0.968 → 0.864	0.983 → 0.947	0.986 → 0.929	–
5 conspiracies	0.965 → 0.752	0.975 → 0.496	0.996 → 0.981	–
25 conspiracies	0.966 → 0.836	0.986 → 0.955	0.997 → 0.985	n/e
AGI	0.958 → 0.690	0.983 → 0.964	0.996 → 0.974	–
Tech	0.897 → 0.586	0.995 → 0.991	0.995 → 0.991	–

Judge collapse index (higher means more collapse)				
Condition	GPT-5 1 h	Gemini Flash Lite 1 h	GPT-5 5 h	Gemini Flash Lite 5 h
Empty	4.075 → 4.319	3.771 → 3.838	4.290 → 4.367	–
1 conspiracy	4.233 → 4.306	4.189 → 4.492	4.478 → 4.636	–
5 conspiracies	4.354 → 4.742	4.316 → 4.894	4.105 → 4.393	–
25 conspiracies	4.434 → 4.642	3.995 → 4.333	4.294 → 4.436	n/e
AGI	4.299 → 4.546	3.910 → 4.175	4.159 → 4.533	–
Tech	4.465 → 4.849	4.131 → 4.327	4.209 → 4.354	–

Table C3: Each cell shows first quarter → final quarter within the same run. Distinct-5 uses the same post count in both windows. One- and five-hour change magnitudes are not directly comparable. “–” means no run. “n/e” means not estimable; the two Gemini Flash Lite five-hour attempts contain three and zero non-seed posts.