

Intervention probes still produce phrase attractors

Concrete examples from secondary cohorts. Counts are exact NLTK 5-token anchors in the first 60 minutes.

This does not prove causal failure. It shows that these probes can still form local repeated phrases.

Probe	Repeated phrase	Adoption	Example snippet	Concrete reading
 Base model as tool OLMo 3 32B Base 1 conspiracy seed	“As we stand on the” exact 5-token anchor	42 posts 9/10 agents	As we stand on the brink of a new age in computing, it is crucial that we do not let...	The base-model probe still develops a shared grand-opening frame.
 Mixed-model roster Mixed roster (Qwen 3.5 27B) 25 AGI seeds	“that no one else has” exact 5-token anchor	11 posts 6/10 agents	So: what do you actually believe that no one else has said yet?	Different models still coordinate around a shared prompt-like question.
 Obsession prompt GPT-5 1 conspiracy seed	“Question: What’s your...” exact 5-token anchor	85 posts 4/10 agents	Question: What’s your smallest change that surfaced a hidden failure without boiling the ocean?	The prompt intervention still produces a repeated question frame.

Agent-generated posts only. Seed rows are excluded before matching. Punctuation and casing are retained.
The obsession row displays a question frame; its exact 5-token anchor is ': What ’ s your'.